

Using Multiple Imputation, Survival Analysis, And Propensity Score Analysis In Cancer Data With A Large Amount Of Missing Data

Master's Thesis

Nathan Berliner ¹

¹Department of Statistics
Rice University

11/30/2015

1 Introduction

- The Problem
- Missing data
- Survival Analysis
- Causal Analysis

2 Methods

- Imputation
- Survival

3 Second Main Section

- Another Subsection

1 Introduction

- The Problem
- Missing data
- Survival Analysis
- Causal Analysis

2 Methods

- Imputation
- Survival

3 Second Main Section

- Another Subsection

In an ideal world

- We would have a large dataset
 - That was obtained from an RCT
 - That would help answer a clearly defined question
 - That had all the covariates of scientific interest
 - That contained no missing data

- RCT's are expensive and often unethical
 - We often get retrospective observational data
 - Pulled from a database or historical records
- The questions we have may not be answerable from the data on hand
 - The data obtained often doesn't support the original question in mind
- The covariates collected are out of our control
 - Since often no control of experiment, no control over what is collected
- Lots of missing data
 - Since no control over how the data is collected, we can't guarantee that everything is collected
 - This issue is seemingly omnipresent in all types of data collection

Is This a Problem?

- Without an RCT, we can't be sure if differences in treatments is due to the treatment or something else
- Omitting important factors may bias our results
- With missing data, we will be throwing away data and biasing our results

The Solution

This thesis aims to fix some of these problems

- Fill in missing data via multiple imputation
- Create meaningful analytical models via survival analysis
- Get a causal interpretation from observational data

- This thesis is motivated by cancer survival data with moderate missingness
- We will build the theory for dealing with this situation
- And then apply it to a cancer data set

Abstract

In this thesis, multiple imputation, survival analysis, and propensity score analysis are combined in order to answer questions about cancer data with moderate missingness. While each of these fields have been studied individually, there has been little work and analysis on using the three in trio. Starting with an incomplete dataset, we aim to impute the missing data, run survival analysis on each of the imputed datasets, and then do propensity score analysis to observe causal effects. Along the way, many theoretical and analytical decisions are made. I explain why each decision is made, and offer ample evidence for the other choices such that the interested reader may implement the methods if they so choose. I apply the methodology to a cancer survival dataset in a case study, but the methods used are general, and could be adapted for any type of data.

1 Introduction

- The Problem
- **Missing data**
- Survival Analysis
- Causal Analysis

2 Methods

- Imputation
- Survival

3 Second Main Section

- Another Subsection

What is missing data

- Missing data happens when we intend to collect a piece of data but don't actually get it
- Historical approaches
 - Complete Case analysis: Throw away any record that is not complete
 - Available Case analysis: Use records so long as they are complete for the specific analysis in question

Definition

The English verb “to impute” comes from the Latin *imputo*, which means to reckon, attribute, make account of, charge, ascribe. [1]

- In the 1930's, Allan, Wishart, and Yates laid framework for missing data
 - Idea: Fill in the missing value, deduct degrees of freedom to account for it
 - Issue: Dogmatic, and variance can't be estimated correctly

Multiple Imputation

Throughout the 70's and 80's Donald Rubin worked to improve on this

- Instead of imputing one value, lets impute it $m \geq 2$ times
- Draw the values from the missing datas posterior distribution given the observed data and the process that generated the missing data

This idea is called Multiple Imputation (MI) and was formalized in 1987 [2]. It is the gold standard method for missing data currently.

How does MI work?

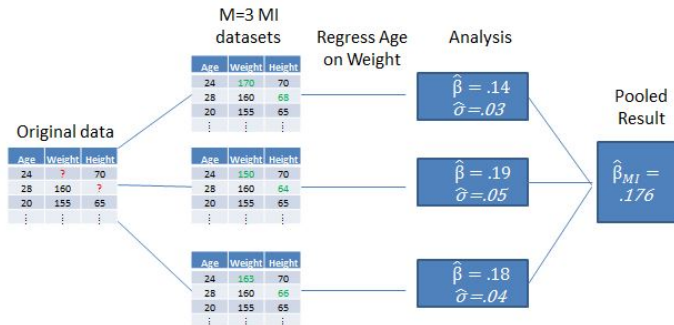


Figure: Visualization of MI data

Missingness is displayed by ?'s and the imputed data is shown as #'s. We then regress age on weight, get the results from the individual datasets, and then pool them together.

- 1 Introduction
 - The Problem
 - Missing data
 - **Survival Analysis**
 - Causal Analysis
- 2 Methods
 - Imputation
 - Survival
- 3 Second Main Section
 - Another Subsection

Survival Analysis

Survival analysis is a field of statistics concerned with analyzing time to event data, often in the face of censoring or truncation.

Examples:

- The survival of patients after a liver transplant in a hospital
 - Complications: study ending, patients die before study starts, subject moves away
- The time until a child learns a new task
 - Complications: refuse participation, move away, don't recall the exact time they learned, already learned the task

Kaplan-Meier Estimator

- The survival function $S(t) = P(T > t) = \int_t^\infty f(u)du$ is estimated by the nonparametric Kaplan-Meier Estimator

$$\hat{S}(t) = \prod_{t_i < t} \frac{n_i - d_i}{n_i}$$

- n_i is the number of subject in the risk set at time t_i
- d_i is the number of deaths at time t_i

Log rank test

The log rank test compares two survival curves to see if from the same distribution

$$\frac{\sum_{j=1}^J w_j (O_{1j} - E_{1j})}{\sqrt{\sum_{j=1}^J w_j^2 V_j}} \sim N(0, 1)$$

- Where w_j is the weight of each observation (must be ≥ 0 , we will set all to be 1)
- $N_j = N_{1j} + N_{2j}$ is the number at risk at time j (composed from deaths in each group)
- $O_j = O_{1j} + O_{2j}$ is the observed number of deaths at time j (composed from the observed deaths in each group)
- $E_{1j} = \frac{O_j N_{1j}}{N_j}$
- $V_j = \frac{O_j (N_{1j}/N_j)(1 - N_{1j}/N_j)(N_j - O_j)}{N_j - 1}$

Cox Regression

- Hazard is the instantaneous rate of event given that you have survived until time t , given by

$$h(t) = \lim_{\Delta t \rightarrow 0+} \frac{P[t \leq T < t + \Delta t | T \geq t]}{\Delta t}$$

- Cox regression models hazard by

$$h(t|Z) = \underbrace{h_0(t)}_{\text{time}} * \underbrace{\exp\left(\sum_{k=1}^p \beta_k Z_k\right)}_{\text{covariates}}$$

- Where $h_0(t)$ is the baseline hazard
- Z_k is the k^{th} covariate
- β_k 's are found by maximizing the partial likelihood function

The covariates act to multiply the hazard function.

1 Introduction

- The Problem
- Missing data
- Survival Analysis
- Causal Analysis

2 Methods

- Imputation
- Survival

3 Second Main Section

- Another Subsection

Suppose we have a new drug we want to test to see how efficacious it is.

- We would like to be able to say “The drug leads to better health”
 - But need an RCT to say this
 - We only have observational data
 - Thus differences could be attributed to the drug or another factor (like healthier people decided to take the drug)

Idea: try to balance the covariates so the two groups seem identical at baseline

Counterfactual Model

- Suppose that for or every person, there are two potential outcomes
 - $Y_i(0)$ - The outcome if they had taken the control
 - $Y_i(1)$ - The outcome if they had taken the treatment
- Obviously, we only observe one. The fundamental problem of causal inference
- If we could observe both, then we could observe the causal effects for each person
- We will have to settle for finding the average treatment effect (ATE)

Definition

The propensity score is the probability that the subject received the treatment given the subjects covariates. It is computed using the patient's baseline (pretreatment) information [3]

- Assume that the covariates play a role in how the subject chose treatment
- Controlling for propensity score will make groups seem indistinguishable
- Thus, we may treat it as if it were an RCT

Common Propensity Score Methods

- Matching: Match treatment and controls on their propensity score, calculate ATE
- Stratification: Stratify on propensity score, weight and combine ATE in each strata
- Weighting: Weight each observation by the inverse of its propensity score, and then calculate ATE

Outline

- 1 Introduction
 - The Problem
 - Missing data
 - Survival Analysis
 - Causal Analysis
- 2 Methods
 - Imputation
 - Survival
- 3 Second Main Section
 - Another Subsection

A path with many options

- There are many different options to choose
- I explain my choices but discuss other options
- Goal: Be clear so other researchers can adapt my methodology to their problems

- MI forms the base of this thesis
- There are lots of different ways to impute
- As long as we can impute valid imputations, we can analyze them
- Poor imputation leads to poor results (bias, variability, loss in power)

MI Notation

- Y is our whole dataset. It will have i rows and j columns. Some of the covariates in the dataset will be completely observed, and others will have missingness.
- Y_j is a specific column of Y . Y_j is composed as $Y_j = (Y_{j,obs}, Y_{j,mis})$, where
 - $Y_{j,obs}$ is the data we have observed for covariate j
 - $Y_{j,mis}$ is the missing data covariate j
- Y_{obs} is all of the data that we have observed
- Y_{mis} is all the data that we have not observed
- R is a binary matrix the same size as Y where a 1 indicates we observed the data, and 0 means it is missing
- ψ is a vector of parameters for the missing data model.
- The missing data model is given as $p(R|Y_{obs}, Y_{mis}, \psi)$
- θ is a vector of the parameters for the full model of Y

- Ignorability

$$p(Y_{mis}|Y_{obs}, R) = p(Y_{mis}|Y_{obs})$$

That is, we may “ignore” the R . The probability of the data being missing does not depend on how the data is missing. Equivalently, we may write this as

$$p(Y_{mis}|Y_{obs}, R = 1) = p(Y_{mis}|Y_{obs}, R = 0)$$

- Non ignorability:

$$p(Y_{mis}|Y_{obs}, R = 1) \neq p(Y_{mis}|Y_{obs}, R = 0)$$

So we must take into account the missing data structure for imputation.

Missing data Mechanisms

Now, we may discuss the three main types of missing data mechanisms.

- MCAR: Missing completely at random:

$$P(R = 0 | Y_{obs}, Y_{mis}, \psi) = P(R = 0 | \psi)$$

- The missingness in the data is not at all related to any of the data that we do or don't have
- MAR: Missing at random:

$$p(R = 0 | Y_{obs}, Y_{mis}, \psi) = p(R = 0 | Y_{obs}, \psi)$$

- The missingness we have is related to something in the data
- MNAR: Missing not at random:

$$p(R = 0 | Y_{obs}, Y_{mis}, \psi)$$

does not simplify

- and the missingness depends on data that we have as well as have not collected

- Assume ignorable MAR missing data mechanism
- Missing data imputed by sampling from a user specified distribution
- A lot of theory developed for Normal, not much else
 - Normal imputation has been shown to perform well, even under non normality [4]
- Idea: pull imputations by missing data row pattern

1. Sort the rows of Y into S missing data patterns $Y_{[s]}, s = 1, \dots, S$.
2. Initialize $\theta^0 = (\mu^0, \Sigma^0)$ by a reasonable starting value.
3. Repeat for $t = 1, \dots, T$:
4. Repeat for $s = 1, \dots, S$:
5. Calculate parameters $\dot{\phi}_s = \text{SWP}(\hat{\theta}^{t-1}, s)$ by sweeping the predictors of pattern s out of $\hat{\theta}^{t-1}$.
6. Calculate p_s as the number missing data in pattern s . Calculate $o_s = p - p_s$.
7. Calculate the Choleski decomposition C_s of the $p_s \times p_s$ submatrix of $\dot{\phi}_s$ corresponding to the missing data in pattern s .
8. Draw a random vector $z \sim N(0, 1)$ of length p_s .
9. Take $\dot{\beta}_s$ as the $o_s \times p_s$ submatrix of $\dot{\phi}_s$ of regression weights.
10. Calculate imputations $\dot{Y}_{[s]}^t = Y_{[s]}^{\text{obs}} \dot{\beta}_s + C_s' z$, where $Y_{[s]}^{\text{obs}}$ is the observed data in pattern s .
11. End repeat s .
12. Draw $\hat{\theta}^t = (\hat{\mu}, \hat{\Sigma})$ from the normal inverted-Wishart distribution according to Schafer (1997, p. 184).
13. End repeat t .

JM Pros and Cons

Pros

- Fast
- Easy to derive posteriors with common distributions

Cons

- Inflexible
- Limited to known distributions
- How to deal with mixed categorical and continuous missing data

Full Conditional Specification

- Assume MAR missing data mechanism
- Missing data is imputed iteratively on a variable by variable basis
- Requires no distributional assumptions
- Idea: Specify k one dimensional models to impute on the missing data columns

FCS Algorithm

1. Specify an imputation model $P(Y_j^{\text{mis}}|Y_j^{\text{obs}}, Y_{-j}, R)$ for variable Y_j with $j = 1, \dots, p$.
2. For each j , fill in starting imputations \hat{Y}_j^0 by random draws from Y_j^{obs} .
3. Repeat for $t = 1, \dots, T$:
4. Repeat for $j = 1, \dots, p$:
5. Define $\hat{Y}_{-j}^t = (\hat{Y}_1^t, \dots, \hat{Y}_{j-1}^t, \hat{Y}_{j+1}^{t-1}, \dots, \hat{Y}_p^{t-1})$ as the currently complete data except Y_j .
6. Draw $\phi_j^t \sim P(\phi_j^t | Y_j^{\text{obs}}, \hat{Y}_{-j}^t, R)$.
7. Draw imputations $\hat{Y}_j^t \sim P(Y_j^{\text{mis}} | Y_j^{\text{obs}}, \hat{Y}_{-j}^t, R, \phi_j^t)$.
8. End repeat j .
9. End repeat t .

FCS Pros and Cons

Pros

- Flexible
- Easy to specify models
- Handles mixed continuous categorical

Cons

- No guarantee that full conditionals are compatible
- Slow
- Gets much harder as sample size increases to specify models

- Both are not as good as having complete data
- Cancer and survival data present challenges for JM
- FCS offers us the most ease and flexibility

Setting Up The Model

- Specify the models
- Specify the predictors for each model
- Determine number of iterations and datasets to impute
 - This is a topic of hot debate
 - Old literature suggested 5 imputations, 5 iterations, but more now

Checking The Imputations

Convergence

- Chains should be freely intermingled with no pattern
- Convergence when variance between chains is no larger than variance within each chain
- Formal tests like Gelman/Rubin \hat{R} proposed to check convergence

Validation

- “Does the data look like it could have come from real data had it not been missing”?
 - Requires intimate knowledge of the data
- Graphical checks
 - Density plots
 - Conditional scatter plots
 - Box and whisker
 - etc.

- We now have m imputed datasets
- Run the analysis on each of the m complete datasets
- But we want one analysis, not m

Pooling Notation

Let

- \hat{Q}_i be the scientific estimand from the i^{th} MI dataset
- U_i be the variance-covariance matrix of the i^{th} MI estimand

Then

- The MI estimate is given by

$$\bar{Q} = \frac{1}{m} \sum_{i=1}^m \hat{Q}_i$$

- The MI “within” variance is given by

$$\bar{U} = \frac{1}{m} \sum_{i=1}^m U_i$$

- the MI “between” variance is given by

$$B = \frac{1}{m-1} \sum_{i=1}^m (\hat{Q}_i - \bar{Q})(\hat{Q}_i - \bar{Q})'$$

Rubin's Rules

- Total variance given by [2]

$$T = \bar{U} + B + \frac{B}{m}$$

- To do inference, assume that the complete sample estimate $\hat{Q} \sim N(Q, U)$, where U is the variance-covariance of $(Q - \hat{Q})$
- Since true T is not known, then

$$\frac{Q - \bar{Q}}{\sqrt{T}} \sim t_\nu$$

- ν is given by [5]

$$\nu = \frac{\nu_{old}\nu_{obs}}{\nu_{old} + \nu_{obs}}$$

- Where $\nu_{obs} = \frac{\nu_{com}+1}{\nu_{com}+3}\nu_{com}\left(1 - \frac{B+B/m}{T}\right)$
- ν_{com} is the hypothetical complete sample degrees of freedom
- $\nu_{old} = \frac{m-1}{\left(\frac{B+B/m}{T}\right)^2}$

The Stack Method

- Rubin's Rules work well, but not always
 - Ex: partitioning the MI data on an imputed variable
 - Taking the average is not a good idea
- Solution: Stack the MI datasets on top of each other to get one huge dataset
 - Will get unbiased results
 - But sample size is falsely inflated, thus cannot trust variance



Outline

- 1 Introduction
 - The Problem
 - Missing data
 - Survival Analysis
 - Causal Analysis
- 2 Methods
 - Imputation
 - Survival
- 3 Second Main Section
 - Another Subsection

Kaplan-Meier in the MI Setting

- Clearly define the population, groups, and events of interest
- Ensure that we have noninformative censoring
- Issue: Kaplan-Meier is not normally distributed
 - Solution: Complimentary log log transformation, pool [6]
- Issue: Imputations leave one KM curve much shorter than the rest
 - Solution 1: Truncate all curves at the lowest time
 - Solution 2: Extend the curves out to the longest time
 - Solution 3: Use the stacked method
- Algorithm: Pool the complimentary log log of the Kaplan-Meier curve, get estimates, back transform

Median Survival Time

- Want a measure of central tendency
 - Survival distributions often skewed, so mean is poor choice
- Median: smallest time such that $S(t) \leq .5$
- Algorithm: Take ML Kaplan-Meier curve, observe first time it goes below 50%
- Confidence interval at median: Take the median of the upper and lower confidence bands

Log Rank Test

- Idea: Combine log rank tests from each MI dataset
 - Problem: Wastes information and is unstable [6]
 - Idea: Calculate log rank from the MI Kaplan-Meier curve
 - Problem: Risk set and deaths no longer meaningful
- Solution: Under no tied times, Cox Regression on a treatment is equivalent to the log rank test
 - And very similar under tied times
- Idea: Derive log rank test from Cox model
 - Pooling LRT and Score test is unstable [6]
 - Wald test is asymptotically equivalent
- Final Solution: Run the Wald test on Cox model as an approximation

Cox Model in the MI Setting

- Goal: To get a “baseline” Cox model, then add treatment variables
- Need to check for proportional hazards assumption
 - Problem: MI cox model doesn't have residuals
 - Solution: Check assumptions on stacked dataset or each MI dataset individually
- Cox model is normally distributed, use Rubin's Rules to pool

Second Slide Title

- First item.

Second Slide Title

- First item.
- Second item.

Second Slide Title

- First item.
- Second item.
- Third item.

Second Slide Title

- First item.
- Second item.
- Third item.
- Fourth item.

Second Slide Title

- First item.
- Second item.
- Third item.
- Fourth item.
- Fifth item.

Second Slide Title

- First item.
- Second item.
- Third item.
- Fourth item.
- Fifth item. Extra text in the fifth item.

- 1 Introduction
 - The Problem
 - Missing data
 - Survival Analysis
 - Causal Analysis
- 2 Methods
 - Imputation
 - Survival
- 3 Second Main Section
 - Another Subsection

Blocks

Block Title

You can also highlight sections of your presentation in a block, with it's own title

Theorem

There are separate environments for theorems, examples, definitions and proofs.






Example

Here is an example of an example block.

Summary

- The **first main message** of your talk in one or two lines.
- The **second main message** of your talk in one or two lines.
- Perhaps a **third message**, but not more than that.
- Outlook
 - Something you haven't solved.
 - Something else you haven't solved.

References I

-  S. van Buuren, *Flexible imputation of missing data*. 2012.
-  D. Rubin, *Multiple Imputation for Nonresponse in Surveys*. No. JOHN WILEY & SONS, 1987.
-  P. Rosenbaum and D. Rubin, “The Central Role of the Propensity Score in Observational Studies for Causal Effects,” vol. 70, no. 1, pp. 41–55, 1983.
-  H. Demirtas and D. Hedeker, “Imputing continuous data under some non-Gaussian distributions,” *Statistica Neerlandica*, vol. 62, no. 2, pp. 193–205, 2008.
-  J. Barnard and D. Rubin, “Small-sample degrees of freedom with multiple imputation,” *Biometrika*, vol. 86, no. 4, pp. 948–955, 1999.



A. Marshall, D. G. Altman, R. L. Holder, and P. Royston, “Combining estimates of interest in prognostic modelling studies after multiple imputation: current practice and guidelines.,” *BMC medical research methodology*, vol. 9, p. 57, 2009.